

Stock Price Prediction Using Twitter

Amal Duriseti
904331855

Jonathan Hurwitz
804258351

Fred Xu
004255573

Jerry Chen
804449266

aduriseti@gmail.com

jdhurwitz@ucla.edu

xuzeyuanfred@ucla.edu hostmissile@hotmail.com

Abstract

This project aims to predict buy or sell indicators for several large S&P500 companies based on Twitter data. The first iteration of our project aimed to predict mood based on tweets, but we pivoted and leveraged learnings from the previously done sentiment analysis work in order to build our stock predictor. Several machine learning models have been tested and compared.

1 Introduction

The former goal was to create an aggregate mood detector and predictor by training machine learning models to perform sentiment analysis on tweets. This task implicitly includes building some method of obtaining more data, such as a data crawler. While our models showed reasonably good performance when trained on the twitter sentiment analysis training corpus¹, we realized that some sort of unsupervised learning would have to be done in order to label future data obtained via the crawler. This motivated us to explore other ways in which we could leverage our previous work but for a task with an easier means of labeling data.

Using sentiment to predict buy or sell indicators seemed like a natural progression with an easier way of labeling data. The labeling scheme we chose was to associate each tweet with the corresponding stock price on the day it was tweeted. The crawler pulls tweets and labels each tweet by doing a yahoo finance price lookup. This way, we are able to build an arbitrarily large dataset of labeled tweets in order to better train our models.

¹The sentiment analysis corpus contains 1.5M tweets and can be downloaded at <http://thinknook.com/twitter-sentiment-analysis-training-corpus->

There is no need for unsupervised learning such as clustering in order to assign class labels.

All code for this project can be found here: https://github.com/aduriseti/cs145_project

2 Data Crawler

The data crawler performs tweet lookup, preprocessing, and labeling for an arbitrary number of stock tickers or query strings. It is assumed that data for each query will be pulled for each day within the specified date range.

Within the scope of this project, the queries were all stock tickers for several large companies: AAPL, AMZN, SNE, GOOG, NVDA, INTC, HPQ, MSFT, IBM, and TXN. Each of these tickers is given their own thread to perform the lookup, filtering, and labeling. The initial intent was to perform a live lookup on Yahoo finance in order to obtain stock prices for labeling, but there were issues with Yahoo finance historical data. After writing a function to pull data using BeautifulSoup, we decided it would be best to just download local csv files from Yahoo containing the data for each ticker over the desired time range (from min to max date).

For each day in the date range, the crawler queries twitter and limits the number of tweets pulled via the "max_tweets" parameter. Tweets are processed in batches corresponding to each day. Tweets for a specific ticker that were posted on a day in which there is no associated stock price are removed from the set. Duplicates are also removed. Stock price labels are 5-D and consist of the open, high, low, close, and adj close. The intent was to provide the maximum amount of information for us to use in learning algorithm evaluation. Following a successful tweet:label pairing, the crawler writes this to a file along with the date

of the tweet and continues.

3 Sentiment Analysis

We trained a variety of classifiers on the Twitter sentiment analysis training corpus. These included random forest, multi-layer perceptron (MLP), and quadratic discriminant analysis. In addition, we tested the RNN-based long short-term memory (LSTM) model. SVM was avoided due to long training time. Metrics used to judge performance include accuracy, precision, recall, and F1 score.

The embedding scheme involved using word2vec and a mean aggregation method. The word2vec vectors are trained on the set of aggregated tweets created by the crawler. For each word within a tweet, we check to see if it is in the trained word2vec model. If so, the word2vec vector corresponding to that word gets added to a running sum for the tweet. If a word does not appear in the word2vec model, then it is zero padded. After all the words have been checked, this summed vector is divided element-wise by the number of words in the tweet in order to obtain a mean vector. This is a form of summarization.

The word2vec model is using dimension 100 vectors, and the resulting summarized tweet vectors are also of dimension 100.

3.1 Random Forest

Random forest provided the initial baseline:

Table 1: word2vec + Random Forest Classifier

Type	Training Set	Test Set
Accuracy	0.684591533095	0.673612152101
Precision	0.681912795793	0.67805020432
Recall	0.82712112152	0.821646476258
F1 score	0.750909267723	0.73868361183




Figure 1: **Figure 1: Random forest classifier ROC with AUC Gain=0.23**




Figure 2: **Figure 2: Random forest classifier precision vs. recall with AUC Gain=0.22**

3.2 Multi-Layer Perceptron

MLP yielded better accuracy on both training and test sets compared to random forest:

Table 2: word2vec + MultiLayer Perceptron (MLP)

Type	Training Set	Test Set
Accuracy	0.811256993379	0.742028552952
Precision	0.828856719281	0.765130190007
Recall	0.840213932107	0.775637595862
F1 score	0.834496685544	0.77034806483




Figure 3: **Figure 3: MLP classifier ROC with AUC gain=0.32**




Figure 5: **Figure 5: QDA classifier ROC with AUC gain=0.28**




Figure 4: **Figure 4: MLP classifier precision vs. recall with AUC Gain=0.34**




Figure 6: **Figure 6: QDA classifier precision vs. recall with AUC Gain=0.26**

3.3 Quadratic Discriminant Analysis

QDA is a Bayesian classifier. Training and test accuracies were better than the baseline random forest but not as high as the multilayer perceptron classifier.

Table 3: word2vec + Quadratic Discriminant Analysis

Type	Training Set	Test Set
Accuracy	0.719376196853	0.71173456698
Precision	0.769485042382	0.760003838403
Recall	0.720252828854	0.706260032103
F1 score	0.744055433157	0.732146984054

3.4 LSTM

The LSTM has a notion of memory. Its memory cell consists of four main elements: an input gate, a neuron with a path back to itself, a forget gate, and an output gate. The loop back allows the state of a memory cell to remain constant through time. The forget gate modulates the loop back and can allow the cell to delete or forget its previous state. LSTMs do not suffer from the vanishing or exploding gradient problem when long sequences are processed. It is an effective model for data that exhibits long-chain dependencies.

3.4.1 LSTM Architecture




Figure 7: **LSTM architecture.**

The figure above shows the basic LSTM architecture, including the notion of time ($t-1$ is introduced) and the gates previously discussed. Corresponding functions are shown below.

$$\begin{aligned}
 z_t &= \sigma(W_z \cdot [h_{t-1}, x_t]) \\
 r_t &= \sigma(W_r \cdot [h_{t-1}, x_t]) \\
 \tilde{h}_t &= \tanh(W \cdot [r_t * h_{t-1}, x_t]) \\
 h_t &= (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t
 \end{aligned}$$

Figure 8: **LSTM functions.**

Our specific network's architecture consists of three layers. Words will be fed into the word embedding layer, which will map words to word embeddings. The resulting word vectors are passed into LSTM cells, whose output will be passed into a dense layer with softmax activation function, and makes a prediction.

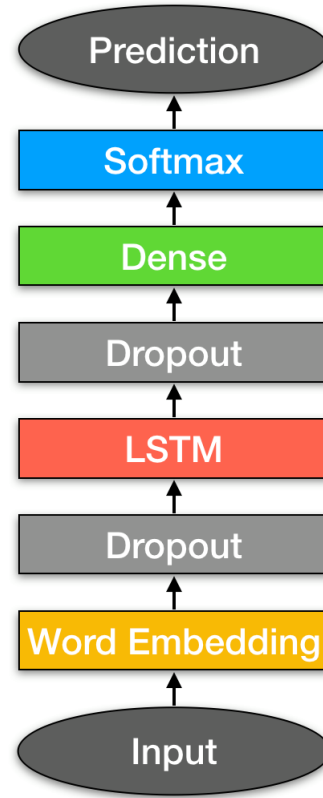


Figure 9: **LSTM implementation architecture.**

3.4.2 LSTM Implementation

Our LSTM model is implemented with Keras, running on top of TensorFlow. Keras is a high-level neural network API, which allowed us to quickly implement the model by simply stacking layers together:

```

model = Sequential()
model.add(Embedding(...))
model.add(Dropout(...))
model.add(LSTM(...))
model.add(Dense(..., activation='softmax'))

```

Our work largely follows Peter Nagy's implementation(?), but we modified the network and tuned hyperparameters to suit our needs. Training the model took two hours to finish.

3.4.3 LSTM Results

The model achieves an F1-score of 0.81 and an Accuracy of 0.81 on test datasets, outperforming the first classifier's F1-score of 0.78 and Accuracy of 0.75.

Metrics	Value
Accuracy	0.810448419184
Precision	0.803874201171
Recall	0.822235970902
F1-score	0.812951416989

4 Stock Prediction

5 Hyperparameter Tuning

The hyperparameter tuning targets the hyperparameters of the word2vec model. Cross-validation was used on the original dataset. Since the MLP classifier had the best performance during the sentiment analysis evaluation phase, we chose to use this as the go-to classifier and invariant pairing. The F1-score of the word2vec + MLP classifier is used as a performance metric. Higher F1 score indicates a better performance.

5.1 Testing Methodology

Word2vec turns text messages into word vectors. However, there is no systematic metric for the evaluation of the model alone. Therefore, evaluation of the bundle: Word2vec + MLP, was used to test the performance.

The default method for hyperparameter tuning is to use an exhaustive grid search. However, since a single run of the 5-fold cross validation takes around 50 minutes, and some hyperparameters, such as size and iter, change the amount of work to be done, a brute force method would be too time consuming. Therefore hyperparameter selection and candidate selection was used. In the selection part, several hyperparameters are examined and the relatively important ones are picked. The resulting candidates for the tuning are:

- Size, which determines the dimensionality of the feature vector.
- `wind_size`, the maximum distance between current and predicted words within a sentence, governs the context of words.
- `min_count`, which defines the cut-off threshold for words.
- `iter`, number of iterations over the corpus.

The candidates for size is picked in the range from the default 100 to 150, `wind_size` is picked from 1 to 10, which includes its default value 5. `min_count` and `iter` are similar to the `wind_size`.

This means a 2-D grid search could take around 3000 testing instances, which is impossible to handle. Therefore, 1-D search was performed on each hyperparameter, and a final search in the potential best candidate's neighborhood was conducted.

The final candidates were further reduced based on their the neighbor's frequency in the 1-D search. If one of the two neighbors in the 1-D case has significantly lower score, it is omitted. Furthermore, size values are limited to only multiple of 4 to improve performance.

5.2 Testing Results

The following table summaries the 1-D search and the resulting search.

hyperparameter	optimum F1	best value	candidate
<code>wind_size</code>	0.6286	1	1,4
<code>size</code>	0.657	130, 140	132, 140
<code>min_count</code>	0.6026	7	7,8
<code>iter</code>	0.6101	9	9, 10

The neighborhood search in the end gives the best candidates for hyperparameters: (140,1,7,9), (140,1,8,9), (140,1,8,10), and (140,4,8,10).

6 Discussion

6.1 Future Work

Due to time limitations, we were not able to integrate the trained LSTM model with the rest of our functional blocks. However, our design was pipelined such that blocks (preprocessor, embedding scheme, sentiment analysis, and prediction) can be easily interchanged.