

```
In [1]: import pandas as pd
import numpy as np
```

```
In [2]: with open("sent_corpus.csv", "r") as sent_file:
lines = sent_file.read().split("\n")
```

```
In [4]: %time rows = [line.split(",") for line in lines if line]
%time rows = [row[:3] + [",".join(row[3:])] for row in rows]
# remove document start character
rows[0][0] = rows[0][0][1:]
%time sentDf = pd.DataFrame(rows[1:], columns=rows[0])
sentDf[["ItemID", "Sentiment"]] = sentDf[["ItemID", "Sentiment"]].astype(int)
%time sentDf["SentimentText"] = sentDf["SentimentText"].apply(lambda text: text.split())
```

```
CPU times: user 4.14 s, sys: 699 ms, total: 4.84 s
Wall time: 5.34 s
CPU times: user 3.93 s, sys: 476 ms, total: 4.41 s
Wall time: 4.96 s
CPU times: user 361 ms, sys: 62 ms, total: 423 ms
Wall time: 464 ms
CPU times: user 6.56 s, sys: 2.4 s, total: 8.96 s
Wall time: 9.29 s
```

```
In [5]: import gensim
%time w2vM = gensim.models.Word2Vec(sentDf["SentimentText"])
```

```
CPU times: user 6min 31s, sys: 5.41 s, total: 6min 36s
Wall time: 2min 55s
```

```
In [6]: # number of samples to aggregate
Ns = int(1e5)
```

```
In [7]: %%time
# MEAN AGGREGATION
tvecs = np.array([np.array([w2vM[t] if t in w2vM
                           else np.zeros((100,))
                           for t in twt]).mean(axis=0)
                  for twt in sentDf["SentimentText"][:Ns]])
# SUM AGGREGATION
# tvecs = np.array([np.array([w2vM[t] if t in w2vM
#                             else np.zeros((100,))
#                             for t in twt]).sum(axis=0)
#                  for twt in sentDf["SentimentText"][:Ns]])
```

```
CPU times: user 10 s, sys: 475 ms, total: 10.5 s
Wall time: 13 s
```

```
In [8]: # number of samples to train on
N = int(1e5)
X = tvecs[:N]
y = sentDf["Sentiment"][:N].values
# inverse classifier
# y = 1 - y
```

```
In [9]: # generate test/train split
ratio = 0.8
tidx = np.random.rand(N) < ratio
pidx = ~tidx
```

```
In [11]: import sklearn
from sklearn import ensemble, svm, neural_network, discriminant_analysis
from sklearn.metrics import roc_curve, auc, precision_recall_curve

from matplotlib import pyplot as plt
% matplotlib inline
```

```
In [12]: def roc_auc(clf, X, y):
    probs = clf.predict_proba(X[pidx])
    fpr, tpr, thresholds = roc_curve(y[pidx], probs[:, 1])
    roc_auc = auc(fpr, tpr)
    roc_str = 'ROC (AUC Gain = %0.2f)' % (roc_auc - 0.5)
    plt.plot(fpr, tpr, lw=1, label=roc_str)
    plt.plot([0, 1], [0, 1], label="RAN CLF")
    plt.title(roc_str)
    plt.show()
```

```
In [13]: def prrc_auc(clf, X, y):
    probs = clf.predict_proba(X)
    pr, rc, thresholds = precision_recall_curve(y, probs[:, 1])
    roc_auc = auc(rc, pr)
    roc_str = 'Prec vs Recall (AUC Gain = %0.2f)' % (roc_auc -
np.mean(y))
    plt.plot(rc, pr, lw=1, label=roc_str)
    plt.plot([0, 1], [np.mean(y), np.mean(y)], label="RAN CLF")
    plt.axis([0, 1, 0, 1])
    plt.title(roc_str)
    plt.show()
```

```
In [14]: def evaluate(clf, X, y):
    yhat = clf.predict(X)
    accu = np.mean(yhat == y)
    prec = np.mean(y[yhat == 1])
    recl = np.mean(yhat[y == 1])
    f1 = 2 * prec * recl / (prec + recl)
    print("Accuracy", accu, "Precision", prec, "Recall", recl, "F1", f1)
```

```
In [15]: rf = sklearn.ensemble.RandomForestClassifier()
rf.max_depth = 5
clf = rf
%time clf.fit(X[tidx],y[tidx])

print("TEST")
%time evaluate(clf,X[pidx],y[pidx])

print("TRAIN")
%time evaluate(clf,X[tidx],y[tidx])

print("ROC AUC")
%time roc_auc(clf,X,y)

print("PRECISION/RECALL AUC")
%time prrc_auc(clf,X,y)
```

CPU times: user 5.87 s, sys: 177 ms, total: 6.05 s

Wall time: 6.84 s

TEST

Accuracy 0.668606675621 Precision 0.659383753501 Recall 0.839664704833

F1 0.73868361183

CPU times: user 44.2 ms, sys: 12.4 ms, total: 56.6 ms

Wall time: 68.5 ms

TRAIN


Accuracy 0.679412243263 Precision 0.670487305061 Recall 0.853253182461

F1 0.750909267723

CPU times: user 133 ms, sys: 28.2 ms, total: 161 ms

Wall time: 169 ms


ROC AUC



CPU times: user 249 ms, sys: 22.6 ms, total: 272 ms

Wall time: 301 ms

PRECISION/RECALL AUC



CPU times: user 385 ms, sys: 37.1 ms, total: 422 ms

Wall time: 464 ms

```
In [16]: mlp = sklearn.neural_network.MLPClassifier()
         clf = mlp
         %time clf.fit(X[tidx],y[tidx])

         print("TEST")
         %time evaluate(clf,X[pidx],y[pidx])

         print("TRAIN")
         %time evaluate(clf,X[tidx],y[tidx])

         print("ROC AUC")
         %time roc_auc(clf,X,y)

         print("PRECISION/RECALL AUC")
         %time prrc_auc(clf,X,y)
```

CPU times: user 1min 5s, sys: 5.6 s, total: 1min 10s

Wall time: 1min 9s

TEST

Accuracy 0.742028552952 Precision 0.765130190007 Recall 0.775637595862

F1 0.77034806483

CPU times: user 72.8 ms, sys: 49 ms, total: 122 ms

Wall time: 116 ms

TRAIN


Accuracy 0.811256993379 Precision 0.828856719281 Recall 0.840213932107

F1 0.834496685544

CPU times: user 267 ms, sys: 132 ms, total: 400 ms

Wall time: 349 ms


ROC AUC



CPU times: user 288 ms, sys: 70.6 ms, total: 359 ms

Wall time: 364 ms

PRECISION/RECALL AUC



CPU times: user 530 ms, sys: 135 ms, total: 665 ms

Wall time: 666 ms

```
In [17]: qda = sklearn.discriminant_analysis.QuadraticDiscriminantAnalysis()
         clf = qda
         %time clf.fit(X[tidx],y[tidx])


         print("TEST")
         %time evaluate(clf,X[pidx],y[pidx])

         print("TRAIN")
         %time evaluate(clf,X[tidx],y[tidx])


         print("ROC AUC")
         %time roc_auc(clf,X,y)

         print("PRECISION/RECALL AUC")
         %time prrc_auc(clf,X,y)
```

CPU times: user 2.59 s, sys: 219 ms, total: 2.81 s
Wall time: 2.33 s
TEST
Accuracy 0.71173456698 Precision 0.760003838403 Recall 0.706260032103 F
1 0.732146984054
CPU times: user 118 ms, sys: 90.6 ms, total: 209 ms
Wall time: 216 ms
TRAIN
Accuracy 0.719376196853 Precision 0.769485042382 Recall 0.720252828854
F1 0.744055433157
CPU times: user 437 ms, sys: 250 ms, total: 687 ms
Wall time: 693 ms
ROC AUC



CPU times: user 345 ms, sys: 101 ms, total: 446 ms
Wall time: 469 ms
PRECISION/RECALL AUC



CPU times: user 752 ms, sys: 282 ms, total: 1.03 s
Wall time: 827 ms